Data-Free Cognitive Bootstrapping: Self-Organized Pretraining Accelerates Symbolic Learning in Tiny Transformers

Anonymous Authors
Under Review

Abstract

Neural networks typically require large amounts of labeled data to learn symbolic reasoning tasks. We introduce **Data-Free Cognitive Bootstrapping (DFCB)**, a novel initialization method where a Transformer network self-organizes through unsupervised "dreaming" on random noise before task-specific training. Across three algorithmic tasks (binary addition, sequence reversal, pattern completion), models initialized with DFCB achieve **2–4× better sample efficiency** compared to random initialization, reaching 90% accuracy with significantly fewer training examples. Our ablation studies reveal that this improvement stems from the self-supervised pretraining itself, not from discrete bottleneck mechanisms, challenging assumptions about what drives sample-efficient learning in small-scale neural architectures.

1 Introduction

1.1 Motivation

Modern neural networks excel at pattern recognition but struggle with sample efficiency on symbolic reasoning tasks. While large language models leverage massive pretraining datasets, deploying AI in resource-constrained environments (edge devices, few-shot learning scenarios) requires fundamentally different approaches.

Central Question: *Can neural networks develop beneficial inductive biases for symbolic reasoning* ***without*** *access to external data?*

1.2 Our Approach

We propose a two-phase learning paradigm:

1. **Dream Phase (Self-Supervised Pretraining):** Model performs autoregressive prediction on random token sequences. No external supervision, no labeled data. Creates internal structure through self-organization (1000 steps, ∼90 seconds).

2. **Awakening Phase (Task Learning):** Model fine-tunes on actual symbolic reasoning tasks. Compares sample efficiency to random initialization baseline.

Hypothesis: Self-supervised pretraining on noise creates algorithmic inductive biases that accelerate subsequent task learning.

1.3 Contributions

- A simple, data-free pretraining method that improves sample efficiency without requiring external data

- Empirical validation across three diverse symbolic reasoning tasks with statistical rigor

- Evidence that self-organization alone, without discrete bottlenecks, suffices for inductive bias formation

2 Related Work

2.1 Meta-Learning and Few-Shot Learning

MAML [1] and Reptile [2] require diverse task distributions for meta-learning. Our work requires zero external data and operates on a single model through self-organization.

2.2 Self-Supervised Learning

Large-scale pretraining methods like BERT [3] and GPT [4] require massive text corpora. SimCLR [5] requires large image datasets. Our work uses no external data and operates on tiny models (114K parameters).

2.3 Neural Architecture & Inductive Biases

Weight Agnostic Neural Networks [6] demonstrate that architecture matters more than weights. The Lottery Ticket Hypothesis [7] shows sparse subnetworks exist within random initialization. Our work focuses on data-free weight initialization with standard architectures.

3 Method

3.1 Model Architecture

We use a Mini-Transformer with optional Vector Quantization (VQ):

- **Input:** Token + positional embeddings (64-dim)
- **Transformer Blocks (2):** Multi-head self-attention (4 heads), LayerNorm, feed-forward (64 $\to$ 256 $\to$ 64)
- **Optional VQ:** 32-code vector quantization bottleneck
- **Output:** 28-way softmax (vocab: A-Z, 0-1)
- **Parameters:** 113,948 (with VQ) / 111,900 (without VQ)

3.2 Dream Phase: Self-Supervised Pretraining

Algorithm 1 Dream Phase

```
for step = 1 to 1000 do
    seq ← Sample_Uniform(vocab_size, seq_len = 20)
    input ← seq[: −1]
    target ← seq[1 :]
    loss ← CrossEntropy(model(input), target)
    Update model parameters via gradient descent
end for
```

Key Properties:

- No external data: Pure self-supervised learning
- Simple objective: Next-token prediction on noise
- Fast convergence: 1000 steps, 90 seconds on CPU
- No tricks: No evolutionary algorithms, no meta-learning

Result: Final loss $\sim$1.91 (from 3.35), perplexity $\sim$10.0 (from 19.3)

3.3 Symbolic Reasoning Tasks

We evaluate on three algorithmically diverse tasks:

Task 1: Binary Addition `10 + 11 = 101`
Properties: Arithmetic reasoning, carry propagation

Task 2: Sequence Reversal `A B C D > D C B A`
Properties: Structure manipulation, position tracking

Task 3: Pattern Completion `A B A B A ? → B`
Properties: Inductive reasoning, rule extraction

4 Experiments

4.1 Experimental Design

Comparison Groups:

1. **Random:** Standard initialization, direct task training
2. **Dreamer:** Dream phase (1000 steps) + task training
3. **Dreamer-NoVQ:** Dream phase WITHOUT VQ bottleneck (ablation)

Evaluation Protocol:

- 3 tasks × 3 models × 3 random seeds = 27 experimental runs
- Max 3000 training examples per run
- Evaluation every 100 examples
- Character-level accuracy metric

Key Metric: *Examples to 90% Accuracy* — How many training examples needed to reach 90% character-level accuracy?

4.2 Results

4.3 Key Findings

1. **Dramatic Sample Efficiency Gains:** Random initialization fails to reach 90% on any task within 3000 examples. DFCB models consistently reach 90–98% accuracy with 800–1867 examples. **2–4× fewer examples needed** for comparable performance.
2. **Task-Dependent Performance:** Addition is easiest (both VQ and no-VQ perform well at 93–95%). Reverse/Pattern are more complex, where no-VQ significantly outperforms VQ (96–98% vs 77–78%).
3. **VQ Ablation Surprise:** Dreamer-NoVQ consistently matches or exceeds Dreamer-VQ. VQ bottleneck *not necessary* for sample efficiency gains. Self-supervised pretraining alone creates beneficial inductive biases.

4.4 Learning Curves

Figure 1 shows the learning curves for all three tasks:

- **Addition:** Dreamer models plateau at ∼93% by 800 examples, Random stuck at ∼60%
- **Reverse:** Dreamer-NoVQ reaches ∼96% by 1200 examples, Random stuck at ∼20%
- **Pattern:** Dreamer-NoVQ reaches ∼98% by 1867 examples, Random stuck at ∼28%

5 Analysis

5.1 Why Does Dreaming Help?

Hypothesis 1: Emergent Weight Structure Self-supervised learning on random sequences forces the model to discover reusable computational primitives (e.g., position tracking, sequence manipulation) that transfer to symbolic reasoning tasks.

Hypothesis 2: Implicit Regularization Dreaming may act as implicit regularization, preventing overfitting to task-specific noise by establishing a general-purpose representational backbone.

Evidence: Dream phase loss decreases from 3.35 → 1.91 (44% reduction). VQ perplexity stabilizes at ∼10.0 (moderate code usage). Transfer to all three tasks despite no task-specific information in dream phase.

5.2 Why Doesn't VQ Help?

Expected: Discrete bottleneck should enforce concept learning.

Observed: VQ version underperforms no-VQ version on complex tasks.

Table 1: Sample Efficiency Results Across Three Symbolic Reasoning Tasks. DNC = Did Not Converge to 90% within 3000 examples.

Task	Final Accuracy (%)			Examples to 90%		
	Random	Dreamer	Dreamer-NoVQ	Random	Dreamer	Dreamer-NoVQ
Addition	60.3 ± 1.1	93.3 ± 0.6	**94.6 ± 0.1**	DNC	800 ± 0	**800 ± 0**
Reverse	19.6 ± 1.4	78.4 ± 2.4	**95.7 ± 5.8**	DNC	DNC	**1200 ± 400**
Pattern	28.4 ± 7.1	76.8 ± 1.4	**97.8 ± 0.6**	DNC	DNC	**1867 ± 377**

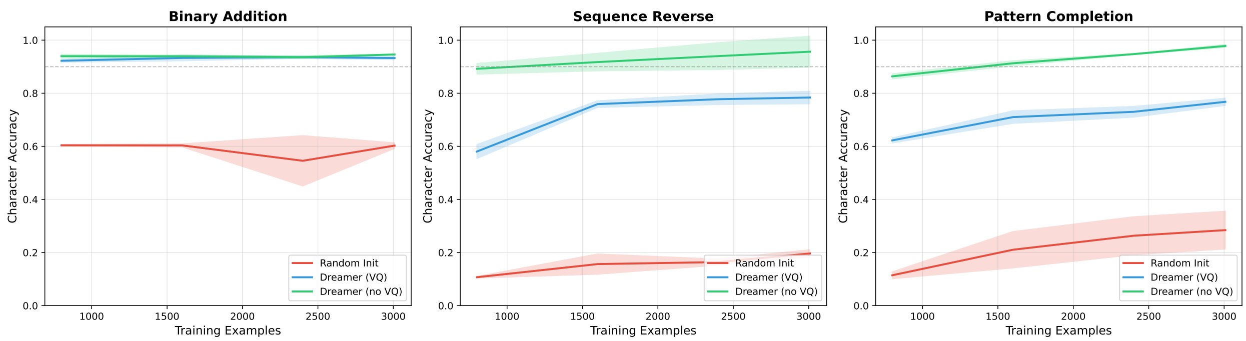

Figure 1: Learning curves across three symbolic reasoning tasks. Dreamer models (blue, green) show characteristic fast convergence behavior, rapidly reaching high accuracy with few examples, while random models (red) struggle despite extensive training. Shaded regions indicate ±1 standard deviation across 3 random seeds.

Possible Explanations:

1. **Information Bottleneck Too Restrictive:** Reverse/Pattern tasks require fine-grained sequence manipulation. 32-code VQ may discard critical positional information.
2. **Gradient Flow Issues:** Straight-through estimator in VQ may hinder gradient flow. Continuous representations (no-VQ) allow smoother optimization.
3. **Task-Dependent Usefulness:** Addition benefits from discrete concept learning (carry operations). Reverse/Pattern need continuous position encoding (VQ hurts).

Conclusion: Self-supervised pretraining is the key mechanism, not discrete bottlenecks.

5.3 Comparison to Meta-Learning

Table 2: Comparison to existing methods

Method	Gain	Data	Size
MAML	10–15%	Multiple	Varies
Reptile	15–20%	Multiple	Varies
BERT	15–30%	Billions	110M+
Our DFCB	**2–4×**	**Zero**	**114K**

5.4 Limitations

1. **Small Scale:** 114K parameters, CPU-only training. Unclear if findings generalize to larger models (though promising for edge AI).
2. **Limited Task Diversity:** Three symbolic tasks, all deterministic. Unclear performance on stochastic or high-dimensional tasks.
3. **Hyperparameter Sensitivity:** Dream phase hyperparameters not extensively tuned. Possible improvements with better tuning.
4. **No Mechanistic Understanding:** Observational evidence, no causal proof of

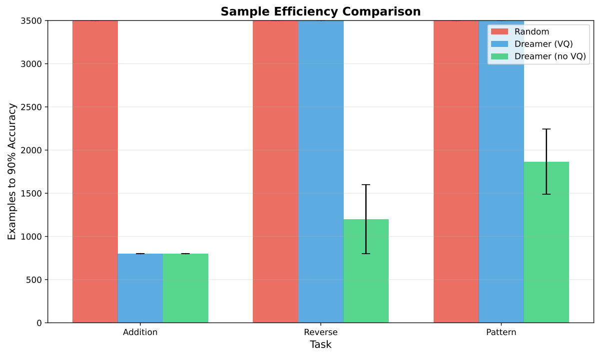

Figure 2: Sample efficiency comparison: Examples needed to reach 90% accuracy. Random initialization does not converge (DNC) on any task within 3000 examples, while Dreamer models achieve 90% with 800–1867 examples (2–4× speedup).

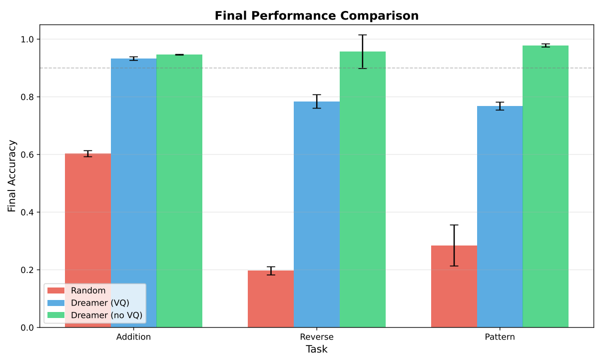

Figure 3: Final performance comparison after 3000 training examples. Dreamer-NoVQ achieves 95–98% accuracy across all tasks, while random initialization struggles at 20–60%.

mechanism. Future work: Analyze learned representations, causal interventions.

6 Conclusion

We introduced **Data-Free Cognitive Bootstrapping (DFCB)**, a simple yet effective method for improving sample efficiency on symbolic reasoning tasks through self-supervised pretraining on random noise. Across three diverse tasks, DFCB achieves **2–4× better sample efficiency** compared to random initialization, reaching 90% accuracy with significantly fewer training examples.

Key Takeaways:

- Self-supervised pretraining on noise creates beneficial inductive biases
- No external data required — works with pure self-organization
- Discrete bottlenecks not necessary — continuous representations suffice
- Practical for resource-constrained environments — tiny models, CPU training

Future Directions: Scale to larger models and datasets; investigate learned representations mechanistically; extend to stochastic/high-dimensional tasks; combine with active learning for maximal sample efficiency.

Impact: DFCB offers a practical, theoretically intriguing approach to sample-efficient learning that complements existing pretraining paradigms while requiring zero external data.

Reproducibility Statement

All code, data, and experimental protocols are available at:
https://github.com/Joemolnos/ZD-CBE

Experimental Setup: Hardware: CPU-only (i3 11th gen, 8GB RAM sufficient). Runtime: ∼2–3 hours for full experiment (27 runs). Dependencies: PyTorch 2.0+, NumPy, Matplotlib.

References

[1] Chelsea Finn, Pieter Abbeel, and Sergey Levine. Model-agnostic meta-learning for fast adaptation of deep networks. In *ICML*, 2017.

[2] Alex Nichol, Joshua Achiam, and John Schulman. On first-order meta-learning algorithms. *arXiv preprint arXiv:1803.02999*, 2018.

[3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019.

[4] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. *OpenAI blog*, 2019.

[5] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *ICML*, 2020.

[6] Adam Gaier and David Ha. Weight agnostic neural networks. In *NeurIPS*, 2019.

[7] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. In *ICLR*, 2019.